

Assignment 3: Wine Sales Poisson Regression

Andrew G. Dunn¹

¹andrew.g.dunn@u.northwestern.edu

Andrew G. Dunn, Northwestern University Predictive Analytics Program

Prepared for PREDICT-411: Generalized Linear Models.

Formatted using the L^AT_EX, references managed using pandoc-citeproc.

Background

This data set contains information on approximately 12,000 commercially available wines. The variables are mostly related to the chemical properties, with others relating to qualitative ratings by individuals. The target variable is the number of sample cases of wine that were purchased by wine distribution companies after sampling. A large wine manufacture is studying the data in order to predict the number of wine cases ordered based upon the wine characteristics. If it is possible to predict the number of cases, the manufacture will be able to adjust their wine offerings with the goal to maximize sales.

We will work towards building Poisson and Negative Binomial models that will predict the target number of cases ordered for each wine.

Exploratory Data Analysis & Data Preparation

We will provide some background into the data after initial examination, then we will:

- Obtain histogram and statistical description for target variable
 - examine mean and variance to check assumption of equality for Poisson or Negative Binomial distribution
 - examine histogram for indication of zero inflation
- Obtain histograms for all continuous variables
 - examine the distributions with transformation in mind during modeling
- Obtain frequency counts for all categorical variables (seeking variability)

Table 1: Data Dictionary

Variable	Description
AcidIndex	Proprietary method of testing total acidity of wine by using a weighted average
Alcohol	Alcohol Content
Chlorides	Chloride content of wine
CitricAcid	Citric Acid Content
Density	Density of Wine
FixedAcidity	Fixed Acidity of Wine
FreeSulfurDioxide	Sulfur Dioxide content of wine
LabelAppeal	Marketing Score indicating the appeal of label design for consumers.
ResidualSugar	Residual Sugar of wine
STARS	Wine rating by a team of experts. 4 Stars = Excellent, 1 Star = Poor
Sulphates	Sulfate conten of wine
TotalSulfurDioxide	Total Sulfur Dioxide of Wine
VolatileAcidity	Volatile Acid content of wine
pH	pH of wine

We're thinking at first that we will model *labelappeal* and *stars* as categorical variables. We will first consider the mean and variance of the target variable:

Table 2: Mean and Variance of the target variable

Mean	Variance
3.0290739	3.7108945

From this we see that we would be in violation of the assumption that mean and variance are equal for the Poisson distribution, however we'd not be in violation for the Negative Binomial distribution which requires the variance to be larger than the mean.

We further examine the histogram of the target variable and notice that it passes tests of normality, but shows obvious signs of being zero inflated. Even though at this stage we would likely restrict our modeling approach based on this observation, we plan to use normal OLS regression, Poisson, and Negative Binomial to examine the performance differences.




Figure 1: Histogram of Target Variable

We then want to examine the variable presence:

Table 3: Variables and Missing Counts

Variable	N	N Missing
AcidIndex	12795	0
Alcohol	12142	653
Chlorides	12157	638
CitricAcid	12795	0
Density	12795	0
FixedAcidity	12795	0
FreeSulfurDioxide	12148	647
LabelAppeal	12795	0
ResidualSugar	12179	616
STARS	9436	3359
Sulphates	11585	1210
TotalSulfurDioxide	12113	682
VolatileAcidity	12795	0
pH	12400	395

We notice that many variables have missing observations within the data set. We will impute these variables with the mean value (in the case of categorical, we will round) and create an indicator for when the data was missing. In this case, with so many, if we find that we want to carry a variable forward into the model, we will be sure to use the missing indicator variable along with it.

Initially looking at the data set we were happy that there was at-least one variable that was going to be an obvious indication of how much the wine was liked (stars), but now we see that there are many missing values for this variable. Due to this we will be interested in the correlation and frequency values for the indicator of imputation for stars.

We examine means statistics for each variable by the count of the target variable and notice there appears to be differences between many variables as we move from zero to eight cases purchased. Not a single variable is sticking out as something that we might not want to use other than the 25% missing observations in stars. We won't eliminate stars because it seems logically like the most valuable variable for indicating desire for purchase.

We examine the two categorical variables (labelappeal, stars) frequency tables given the target variable and notice proportional variation throughout the range.

We examine the histogram of each variable and notice the following:

- each variable passes all tests for normality at the significant level
- many variables have significant center spikes
- stars, after imputation, is left centered

We speculate that the center spike in each of the continuous variables might be caused by a log scaling of the variables. We at this time won't begin scaling because we know it adds complexity to the model interpretation.

We examine the simple Pearson correlation between our imputed variables:

Table 4: Variable Correlation with Target

Variable	Correlation	$\Pr > r $ under $H_0 : \rho = 0$
acidindex	-0.24605	< 0.0001
imp_alcohol	0.06043	< 0.0001
imp_chlorides	-0.03724	< 0.0001
citricacid	0.00868	0.3260
density	-0.03552	< 0.0001
fixedacidity	-0.04901	< 0.0001
imp_freesulfurdioxide	0.04269	< 0.0001
labelappeal	0.35650	< 0.0001
imp_residualsugar	0.01607	0.0691
imp_stars	0.40013	< 0.0001
i_imp_stars	-0.57158	< 0.0001
imp_sulphates	-0.03691	< 0.0001
imp_totalsulfurdioxide	0.05010	< 0.0001
volatileacidity	-0.08879	< 0.0001
imp_ph	-0.00928	0.2939

From this point we're feeling the need to throw out any variable that doesn't have at least an $|0.1|$ level of correlation, without transformation, with our dependent variable. We do see high correlation between the qualitative review variables, *labelappeal* and *stars*, and the dependent variable. Interestingly our highest correlation is with the indicator of imputation for *imp_stars*. Of the variables above, we're going to continue detailed examination into the modeling stage of all those that are significant and above the $|0.1|$ threshold. This leaves us with *acidindex*, *labelappeal*, *imp_stars*, and *i_imp_stars*.

It's quite interesting to be down to one physical measure and two qualitative measures. We're going to try to build the best model we can based on our abilities, however in the conclusion section we will remark on the down-select of variables and what that means to the likely strategy of our patron in this modeling exercise.

In looking through each of the categorical variables *labelappeal*, *imp_stars*, and *i_imp_stars* in a frequency table with the dependent variable we observe what we think is a linear relationship throughout the range.

For Model construction we know ultimately we will be comparing to an OLS regression model, however we will choose to work with Poisson and Negative Binomial models to help us select our parameters. Even though above we communicate that we are down-selecting to a limited range of variables, we spend time examining different combinations of variables and the respective performance to our target.

The Genmod procedures in SAS don't provide us with a method for automatic variable selection, instead we must narrow down by attempting to explore comprehensively what variables make sense to incorporate. This becomes more difficult when we're working with the Zero Inflated variants of the model as we need to produce frequency tables to examine which variables conditionally contribute to the probability that we would observe a zero count in the target variable.

We expect to see *very* similar models from the Poisson and Negative Binomial approaches due to the variance being close to equal with the mean.

Model Construction: Poisson

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \beta_4 X_4 + \epsilon$$

Where:

Table 5: Poisson Model Variables

In Model	In Data
Y is	target
X_1 is	acidindex
X_2 is	labelappeal
X_3 is	imp_stars
X_4 is	i_imp_stars

Table 6: Poisson Model Analysis Of Maximum Likelihood Parameter Estimates

Parameter	Set	Estimate	Wald Chi-Square	Pr > ChiSq
Intercept		1.3752	836.24	<.0001
AcidIndex		-0.0814	328.69	<.0001
LabelAppeal	-2	-0.6958	269.03	<.0001
LabelAppeal	-1	-0.4597	338.98	<.0001
LabelAppeal	0	-0.2702	139.87	<.0001
LabelAppeal	1	-0.1377	35.38	<.0001
LabelAppeal	2	0	.	.
imp_stars	1	-0.5647	682.89	<.0001
imp_stars	2	-0.2431	149.78	<.0001
imp_stars	3	-0.1207	35.77	<.0001
imp_stars	4	0	.	.
i_imp_stars	0	1.0926	3599.71	<.0001
i_imp_stars	1	0	.	.

Table 7: Poisson Model Criteria For Assessing Goodness Of Fit

Criterion	DF	Value	Value/DF
Deviance	1.30E+004	13700.3624	1.0716
Scaled Deviance	1.30E+004	13700.3624	1.0716
Pearson Chi-Square	1.30E+004	11331.5923	0.8863
Scaled Pearson X2	1.30E+004	11331.5923	0.8863
Log Likelihood		8775.9792	
Full Log Likelihood		-22821.192	
AIC (smaller is better)		45664.3841	
AICC (smaller is better)		45664.4047	
BIC (smaller is better)		45746.409	

The exponentiated acidindex coefficient is the multiplicative term used to calculate the estimated target when acidindex increases by 1 unit. In the case of the categorical variables, the exponentiated coefficient is the multiplicative term relative to the base level for each variable. The exponentiated intercept is the baseline

rate, and all other estimates will be relative to it.

The effect of a one unit increase in acidindex is a 8% decrease in the expected number of cases purchased.

Given that labelappeal has a base level of 2 (highest rating), we interpret obtaining a:

- negative 2 rating : 50% decrease in the expected number of cases purchased
- negative 1 rating : 36% decrease in the expected number of cases purchased
- zero rating : 23% decrease in the expected number of cases purchased
- positive 1 rating : 12% decrease in the expected number of cases purchased

Each of which is interpreted as a ‘decrease from from obtaining a 2 rating’.

Given that stars has a base level of 4 (highest rating), we interpret obtaining a:

- 1 rating : 43% decrease in the expected number of cases purchased
- 2 rating : 21% decrease in the expected number of cases purchased
- 3 rating : 11% decrease in the expected number of cases purchased

Each of which is interpreted as a ‘decrease from from obtaining a 4 rating’.

If the wine was given a a star’ed rating, rather than being imputed by us, our model indicates a 98% increase in the expected number of cases to be purchased.

Model Construction: Negative Binomial

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \beta_4 X_4 + \epsilon$$

Where:

Table 8: Negative Binomial Model Variables

In Model	In Data
Y is	target
X_1 is	acidindex
X_2 is	labelappeal
X_3 is	imp_stars
X_4 is	i_imp_stars

Table 9: Negative Binomial Model Analysis Of Maximum Likelihood Parameter Estimates

Parameter	Set	Estimate	Wald Chi-Square	Pr > ChiSq
Intercept		1.3752	836.24	<.0001
AcidIndex		-0.0814	328.69	<.0001
LabelAppeal	-2	-0.6958	269.03	<.0001
LabelAppeal	-1	-0.4597	338.98	<.0001
LabelAppeal	0	-0.2702	139.87	<.0001
LabelAppeal	1	-0.1377	35.38	<.0001
LabelAppeal	2	0	.	.
imp_stars	1	-0.5647	682.89	<.0001
imp_stars	2	-0.2431	149.78	<.0001
imp_stars	3	-0.1207	35.77	<.0001
imp_stars	4	0	.	.
i_imp_stars	0	1.0926	3599.71	<.0001
i_imp_stars	1	0	.	.

Table 10: Negative Binomial Model Criteria For Assessing Goodness Of Fit

Criterion	DF	Value	Value/DF
Deviance	1.30E+004	13700.3624	1.0716
Scaled Deviance	1.30E+004	13700.3624	1.0716
Pearson Chi-Square	1.30E+004	11331.5923	0.8863
Scaled Pearson X2	1.30E+004	11331.5923	0.8863
Log Likelihood		8775.9792	
Full Log Likelihood		-22821.192	
AIC (smaller is better)		45664.3841	
AICC (smaller is better)		45664.4047	
BIC (smaller is better)		45746.409	

The interpretation of results for this model is the same as the model above, the parameters are the very close to the same due to the mean and variance being so close.

We varied inputs for quite a while and didn't find any particular model to be compelling over what we'd achieved with the above parameters.

Model Construction: Zero Inflated Poisson

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \beta_4 X_4 + \epsilon$$

Where:

Table 11: Zero Inflated Poisson Model Variables

In Model	In Data
Y is	target
X_1 is	acidindex
X_2 is	labelappeal
X_3 is	imp_stars
X_4 is	i_imp_stars

Table 12: Zero Inflated Poisson Model Analysis Of Maximum Likelihood Parameter Estimates

Parameter	Set	Estimate	Wald Chi-Square	Pr > ChiSq
Intercept		1.875	1413.66	<.0001
AcidIndex		-0.0223	20.34	<.0001
LabelAppeal	-2	-0.9652	484.26	<.0001
LabelAppeal	-1	-0.5995	533.13	<.0001
LabelAppeal	0	-0.339	206.89	<.0001
LabelAppeal	1	-0.1567	43.46	<.0001
LabelAppeal	2	0	0	0
imp_stars	1	-0.417	327.39	<.0001
imp_stars	2	-0.2012	101.76	<.0001
imp_stars	3	-0.1049	26.98	<.0001
imp_stars	4	0	0	0
i_imp_stars	0	0.1868	90.62	<.0001
i_imp_stars	1	0	0	0

Table 13: Zero Inflated Poisson Model Analysis Of Maximum Likelihood Zero Inflation Parameter Estimates

Parameter	Set	Estimate	Wald Chi-Square	Pr > ChiSq
Intercept		-3.4731	305.03	<.0001
AcidIndex		0.4773	372.21	<.0001
i_imp_stars	0	-3.6189	1550.75	<.0001
i_imp_stars	1	0	0	0

Table 14: Zero Inflated Poisson Model Criteria For Assessing Goodness Of Fit

Criterion	DF	Value	Value/DF
Deviance		41927.6145	
Scaled Deviance		41927.6145	
Pearson Chi-Square	13000	6122.3756	0.479
Scaled Pearson X2	13000	6122.3756	0.479
Log Likelihood		10633.364	
Full Log Likelihood		-20963.8072	
AIC (smaller is better)		41953.6145	
AICC (smaller is better)		41953.6429	
BIC (smaller is better)		42050.553	

The exponentiated acidindex coefficient is the multiplicative term used to calculate the estimated target when acidindex increases by 1 unit. In the case of the categorical variables, the exponentiated coefficient is the multiplicative term relative to the base level for each variable. The exponentiated intercept is the baseline rate, and all other estimates will be relative to it.

The effect of a one unit increase in acidindex is a 2% decrease in the expected number of cases purchased.

Given that labelappeal has a base level of 2 (highest rating), we interpret obtaining a:

- negative 2 rating : 61% decrease in the expected number of cases purchased
- negative 1 rating : 45% decrease in the expected number of cases purchased
- zero rating : 28% decrease in the expected number of cases purchased
- positive 1 rating : 14% decrease in the expected number of cases purchased

Each of which is interpreted as a ‘decrease from from obtaining a 2 rating’.

Given that stars has a base level of 4 (highest rating), we interpret obtaining a:

- 1 rating : 34% decrease in the expected number of cases purchased
- 2 rating : 18% decrease in the expected number of cases purchased
- 3 rating : 10% decrease in the expected number of cases purchased

Each of which is interpreted as a ‘decrease from from obtaining a 4 rating’.

For the Zero inflation, this portion of the model refers to the logistic model (because of our link function) predicting whether or not the number of cases purchased is zero.

The effect of a one unit increase in acidindex is a 61% increase in the odds that this wine would belong to the certain zero group for number of cases purchased. The effect of not being rated, or being omitted from having a star rating, is a 97% increase in the odds that this wine would belong to the certain zero group for number of cases purchased.

Model Construction: Zero Inflated Negative Binomial

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \beta_4 X_4 + \epsilon$$

Where:s that the predicted odds of a student with mathnce and langnce scores of zero being a certain zero are .00137614 (though remember that evaluating mathnce and langnce

Table 15: Zero Inflated Negative Binomial Model Variables

In Model	In Data
Y is	target
X_1 is	acidindex
X_2 is	labelappeal
X_3 is	imp_stars
X_4 is	i_imp_stars

Table 16: Zero Inflated Negative Binomial Model Analysis Of Maximum Likelihood Parameter Estimates

Parameter	Set	Estimate	Wald Chi-Square	Pr > ChiSq
Intercept		1.8705	1403.01	<.0001
AcidIndex		-0.0214	18.71	<.0001
LabelAppeal	-2	-0.9704	487.27	<.0001
LabelAppeal	-1	-0.6029	536.36	<.0001
LabelAppeal	0	-0.3409	207.84	<.0001
LabelAppeal	1	-0.1574	43.55	<.0001
LabelAppeal	2	0	0	0
imp_stars	1	-0.4068	312.88	<.0001
imp_stars	2	-0.1999	99.53	<.0001
imp_stars	3	-0.1046	26.56	<.0001
imp_stars	4	0	0	0
i_imp_stars	0	0.1854	88.92	<.0001
i_imp_stars	1	0	0	0
dispersion		0.0019		

Table 17: Zero Inflated Negative Binomial Model Analysis Of Maximum Likelihood Zero Inflation Parameter Estimates

Parameter	Set	Estimate	Wald Chi-Square	Pr > ChiSq
Intercept		-3.3657	304.21	<.0001
AcidIndex		0.4637	374.47	<.0001
i_imp_stars	0	-3.4689	1757.03	<.0001
i_imp_stars	1	0	0	0

Table 18: Zero Inflated Negative Binomial Model Criteria For Assessing Goodness Of Fit

Criterion	DF	Value	Value/DF
Deviance		41984.4131	

Criterion	DF	Value	Value/DF
Scaled Deviance		41984.4131	
Pearson Chi-Square	13000	6016.3408	0.4707
Scaled Pearson X2	13000	6016.3408	0.4707
Log Likelihood		-20992.2065	
Full Log Likelihood		-20992.2065	
AIC (smaller is better)		42012.4131	
AICC (smaller is better)		42012.4459	
BIC (smaller is better)		42116.8084	

The results of this model are very similar to the model above.

Model Interpretation Comparison

The comparison between model coefficients makes sense, however we would have liked to have an analog to the ROC curve we used in logistic regression (possible, but didn't construct for these models). We were not surprised to see that there was little difference between Poisson and Negative Binomial due to the mean and variance being very close in value. Given that we observed the target value being zero inflated, we chose to continue with the Zero Inflated Negative Binomial Model.

We built comparison methods using a sub-sampling test procedure (not included in analytical codebase) to subjectively evaluate the performance of the difference models.

Conclusion

Logistic and Poisson models have initial complexity when building out interpretation, however once locked in the interpretation of these feels more natural than other techniques we've used in the course. The choice of what to use for a basis of interpretation when incorporating categorical variables can significantly influence the feeling of interpretation.

Without reading too much into the assignment subtext, we wanted to observe that the patron for this study would most likely desire using objectively measurable parameters in their model. Choosing the qualitative/subjective parameters, like we have done, may provide for better model performance, however forces requirements onto future data collection. For instance, the patron may object to the inclusion of the star variable because it requires human sampling of the product in a disposable fashion, where as the labelapproval can be sampled in a non-destructive fashion.

Appendix: Analysis

```
libname four11 '/scs/wtm926/' access=readonly;

Data eda;
    set four11.wine;

proc contents data=eda;

proc univariate data=eda normal;
    var target;
    histogram;

proc means data=eda n nmiss mean var;
    var acidindex alcohol chlorides citricacid density fixedacidity freesulfurdioxide labelappeal residu

proc means data=eda n nmiss mean var;
    var acidindex alcohol chlorides citricacid density fixedacidity freesulfurdioxide labelappeal residu
    class target;

proc means data=eda n nmiss mean var;
    var acidindex alcohol chlorides citricacid density fixedacidity freesulfurdioxide residualsugar sul

proc univariate data=eda normal;
    var acidindex alcohol chlorides citricacid density fixedacidity freesulfurdioxide labelappeal residu
    histogram;

data imp_eda;
    set eda;

    imp_alcohol = alcohol;
    i_imp_alcohol = 0;
    if missing(imp_alcohol) then do;
        imp_alcohol = 10.4892363;
        i_imp_alcohol = 1;
    end;

    imp_chlorides = chlorides;
    i_imp_chlorides = 0;
    if missing(imp_chlorides) then do;
        imp_chlorides = 0.0548225;
        i_imp_chlorides = 1;
    end;

    imp_freesulfurdioxide = freesulfurdioxide;
    i_imp_freesulfurdioxide = 0;
    if missing(imp_freesulfurdioxide) then do;
        imp_freesulfurdioxide = 30.8455713;
        i_imp_freesulfurdioxide = 1;
    end;

    imp_residuaisugar = residuaisugar;
    i_imp_residuaisugar = 0;
    if missing(imp_residuaisugar) then do;
```

```

        imp_residualsugar = 5.4187331;
        i_imp_residualsugar = 1;
    end;

    imp_stars = stars;
    i_imp_stars = 0;
    if missing(imp_stars) then do;
        imp_stars = 2.0;
        i_imp_stars = 1;
    end;

    imp_sulphates = sulphates;
    i_imp_sulphates = 0;
    if missing(imp_sulphates) then do;
        imp_sulphates = 0.5271118;
        i_imp_sulphates = 1;
    end;

    imp_totalsulfurdioxide = totalsulfurdioxide;
    i_imp_totalsulfurdioxide = 0;
    if missing(imp_totalsulfurdioxide) then do;
        imp_totalsulfurdioxide = 120.7142326;
        i_imp_totalsulfurdioxide = 1;
    end;

    imp_ph = ph;
    i_imp_ph = 0;
    if missing(imp_ph) then do;
        imp_ph = 3.2076282;
        i_imp_ph = 1;
    end;

proc means data=imp_eda n nmiss mean var;
    var acidindex imp_alcohol imp_chlorides citricacid density fixedacidity imp_freesulfurdioxide label;

proc freq data=eda;
    tables target*labelappeal;
proc freq data=eda;
    tables target*stars;

proc univariate data=eda normal;
    var acidindex imp_alcohol imp_chlorides citricacid density fixedacidity imp_freesulfurdioxide label;
    histogram;

proc corr data=imp_eda rank plots=all;
    var acidindex imp_alcohol imp_chlorides citricacid density fixedacidity imp_freesulfurdioxide label;
    with target;

proc freq data=imp_eda;
    tables target*i_imp_stars;
proc freq data=imp_eda;
    tables target*imp_stars;
proc freq data=imp_eda;

```

```

tables target*labelappeal;

proc sort data=imp_eda out=zero_check;
  by acidindex;

proc freq data=zero_check;
  table target / plots(only)=freqplot(scale=percent);
  by acidindex;

proc genmod data=imp_eda;
  class labelappeal imp_stars i_imp_stars;
  model target = acidindex labelappeal imp_stars i_imp_stars / link=log dist=poi;
  output out=imp_eda p=pr1;

proc genmod data=imp_eda;
  class labelappeal imp_stars i_imp_stars;
  model target = acidindex labelappeal imp_stars i_imp_stars / link=log dist=nb;
  output out=imp_eda p=nbr1;

proc genmod data=imp_eda;
  class labelappeal imp_stars i_imp_stars;
  model target = acidindex labelappeal imp_stars i_imp_stars / link=log dist=ZIP;
  zeromodel acidindex i_imp_stars / link=logit;
  output out=imp_eda p=zip1;

proc genmod data=imp_eda;
  class labelappeal imp_stars i_imp_stars;
  model target = acidindex labelappeal imp_stars i_imp_stars / link=log dist=ZIP;
  zeromodel acidindex i_imp_stars / link=logit;
  output out=imp_eda p=zip1 pzero=zzip1;

proc genmod data=imp_eda;
  class labelappeal imp_stars i_imp_stars;
  model target = acidindex labelappeal imp_stars i_imp_stars / link=log dist=ZINB;
  zeromodel acidindex i_imp_stars / link=logit;
  output out=imp_eda p=zinb1 pzero=zzinb1;

proc reg data=imp_eda;
  model target = acidindex labelappeal imp_stars i_imp_stars;
  output out=imp_eda p=yhat;

proc genmod data=imp_eda;
  class labelappeal imp_stars i_imp_stars;
  model target = acidindex labelappeal imp_stars i_imp_stars / link=identity dist=normal;
  output out=imp_eda p=ols1;

proc print data=imp_eda (obs=20);
  var target pr1 nbr1 zip1 zinb1 yhat ols1;

run;

```

Appendix: Deployment

```
libname four11 '/scs/wtm926/' access=readonly;

Data testing;
  set four11.wine_test;

data testing_fixed;
  set testing;

  imp_stars = stars;
  i_imp_stars = 0;
  if missing(imp_stars) then do;
    imp_stars = 2.0;
    i_imp_stars = 1;
  end;

data testing_score;
  set testing_fixed;

  TEMP = -3.3657
  + AcidIndex * 0.4637
  + (i_imp_stars in (0)) * -3.4689;

  P_SCORE_ZERO = exp(TEMP) / (1 + exp(TEMP));

  temp = 1.8705
  + AcidIndex * -0.0214
  + (LabelAppeal in (-2)) * -0.9704
  + (LabelAppeal in (-1)) * -0.6029
  + (LabelAppeal in (0)) * -0.3409
  + (LabelAppeal in (1)) * -0.1574
  + (imp_stars in (1)) * -0.4068
  + (imp_stars in (2)) * -0.1999
  + (imp_stars in (3)) * -0.1046
  + (i_imp_stars in (0)) * 0.1854;

  P_SCORE_ZIP_ALL = exp(TEMP);

  P_TARGET = P_SCORE_ZIP_ALL * (1 - P_SCORE_ZERO);
  P_TARGET_ROUND = round(P_TARGET,1);

  keep index target P_TARGET P_TARGET_ROUND;

proc export data=testing_score
  outfile='/sscc/home/a/agd808/sasuser.v94/411/3/out.csv'
  dbms=csv
  replace;
```